

Exchange of Control

John Skaller

December 25, 2016

1 Objects

A coroutine system consists of the following types of objects:

Scheduler A device to hold a set of active fibres and select one to be current.

Channels An object to support synchronisation and data transfer.

Fibres A thread of control which can be suspended and resumed.

Continuations An object representing the future of a coroutine.

1.1 Fibre States

Each fibre is in one of these states:

Running Exactly one fibre (per pthread) is always running.

Active Fibres which are ready to run but not running.

Hungry Fibres suspended waiting for input on a channel.

Blocked Fibres suspended waiting to perform output on a channel.

Dead A fibre for which there is no longer a current continuation.

1.2 Channel States

Each channel is in one of these states:

Empty There are no fibres associated with the channel.

Hungry A set of hungry fibres are waiting for input on the channel.

Blocked A set of blocked fibres are waiting to perform output on the channel.

2 Abstract Operations

Several operations are specific to a fibrated system.

2.1 Spawn

The spawn operation takes as an argument a unit procedure and makes a closure thereof the initial continuation of a new fibre. Of the pair consisting of the currently running fibre (the spawner) and the new fibre (the spawnee) one will have Active state and the other will be Running. It is not specified which of the pair is Running.

2.2 Create channel

A function which creates a channel.

2.3 Read

The read operation takes as an argument a channel.

1. If the channel is Empty, the Running fibre performing the read changes state to Hungry, the channel changes state to Hungry, and the fibre is associated with the channel.

If there are no active fibres, the program terminates, otherwise the scheduler selects an Active fibre and changes its state to Running. It is not specified which active fibre is chosen.

2. If the channel is Hungry, the Running fibre changes state to Hungry, and the fibre is associated with the channel.

If there are no active fibres, the program terminates, otherwise the scheduler selects an Active fibre and changes its state to Running. It is not specified which active fibre is chosen.

3. If the channel is Blocked, one of the associated Blocked fibres is selected, and dissociated from the channel. Of these two fibres, one is changed to state Active and the other to Running. It is not specified which fibre is chosen to be Running.

The value supplied to the write operation of the Blocked fibre will be pass to the Hungry fibre when it transitions to Running state.

2.4 Write

The write operation takes two arguments, a channel and a value to be written.

1. If the channel is Empty, the Running fibre performing the write changes state to Blocked, the channel changes state to Blocked, and the fibre is associated with the channel.

If there are no active fibres, the program terminates, otherwise the scheduler selects an Active fibre and changes its state to Running. It is not specified which active fibre is chosen.

2. If the channel is Blocked, the Running fibre changes state to Blocked, and the fibre is associated with the channel.

If there are no active fibres, the program terminates, otherwise the scheduler selects an Active fibre and changes its state to Running. It is not specified which active fibre is chosen.

3. If the channel is Hungry, one of the associated Hungry fibres is selected, and dissociated from the channel. Of these two fibres, one is changed to state Active and the other to Running. It is not specified which fibre is chosen to be Running.

The value supplied by the write operation of the Blocked fibre will be passed to the Hungry fibre when it transitions to Running state.

2.5 Reachability

The Running, and, each Active fibre and its associated call chain of continuations are deemed to be Reachable.

If a channel is known to reachable fibre, it is also reachable.

Each fibre associated with a reachable channel is reachable.

The transitive closure of the reachability relation consists of a closed, finite, collection of channels and fibres which are reachable.

Unreachable fibres and channels are automatically garbage collected.

2.6 Elimination

Fibres and channels are eliminated when they are no longer reachable.

A fibre may become unreachable in three ways.

2.6.1 Suicide

A fibre for the the initial continuation returns acquires the Dead state. If there are no longer any Active fibres, the program returns, otherwise the scheduler picks one Active fibre and changes its state to Running.

A fibre in the Dead state should eventually become unreachable if it is not unreachable when it becomes Dead.

2.6.2 Starvation

A fibre in the Hungry state may become unreachable, in which case it is marked Dead and is no longer part of the fibration system.

2.6.3 Blockage

A fibre in the Blocked state may become unreachable in which case it is marked Dead and is no longer part of the fibration system.

3 LiveLock

If a fibre is Hungry (or Blocked) on a reachable channel but no future Running fibre will write (or read) that channel, the fibre is said to be livelocked. The fibre will never proceed but it cannot be removed from the system because it is reachable via the channel.

A livelock is considered to transition to a deadlock if the channel becomes unreachable, in which case the fibre will become Dead through Starvation (or Blockage), dissolving the deadlock. In other words, fibres cannot deadlock.

A livelock which persists will hang the program.

4 Fibre Structure

Each fibre consists of a single current continuation. Each continuation may have an associated continuation known as its caller. The initial continuation of a freshly spawned fibre has no caller.

The closure of the caller relation leads to a linear sequence of continuations starting with the current continuation and ending with the initial continuation of a freshly spawned fibre.

The main program consists of an initially Running fibre with a specified initial continuation.

Continuations have the usual operations of a procedure. They may return, call another procedure, spawn new fibres, create channels, and read and write channels, as well as the other usual operations of a procedure in a general purpose programming language.

A continuation is reachable if it is the current continuation of a reachable fibre, or the caller of a reachable continuation.

5 Continuation Structure

5.1 Continuation Data

A continuation has associated with it the following data:

caller Another continuation of the same fibre which is suspended until this continuation returns.

data frame Sometimes called the stack frame, contains local variables the continuation may access.

program counter A location in the program code representing the current point of this continuations execution or suspension

5.2 Continuation operations

The current continuation of a fibre executes a wide range of operations including channel I/O, spawning new fibres, calling a procedure, and returning.

call Calling a procedure creates a new continuation with its program counter set at the procedure entry point, and a fresh data frame. The new continuation becomes the current continuation, the current continuation suspends, and the caller of the new continuation is set to the current continuation. The current continuation program counter is set to the pointer after the call instruction.

return Returning from the current continuation causes the owning fibres current continuation to be set to the current continuations caller, if one exists, or the fibre to be marked Dead if there is no caller. Execution of the suspended caller continues at its program counter.

read/write Channel I/O suspends the current continuation of a fibre until a matching operation from another fibre synchronises with it. A read is matched by a write, and a write is matched by a read.

By the rules of state change, channel I/O should be viewed as performing a peer to peer neutral exchange of control: the current fibre becomes suspended

without losing its position and hands control to another fibre. Later, control is handed back and the fibre continues.

Coroutine based systems, therefore, operate by repeated exchanges of control accompanied by data transfers in a direction independent of the control flow, which sets coroutines aside from functions.

6 Events

Each state transfer of the fibration system may be considered an event. However the key events are

- spawning
- suicide
- entry to a read operation
- return from a read operation
- entry to a write operation
- return from a write operation

I/O synchronisation consists of suspension on entry to a read or write operation, and simultaneously release of suspension, or resumption, on matching write or read.

I/O suspension occurs when a fibre becomes Hungry or Blocked, and resumption when it becomes Running or Active.

Fibrated systems are characterised by a simple rule: events are totally ordered. The order may not be determinate.

7 Control Type

The control type of a coroutine is defined as follows. We assume the coroutine is spawned as a fibre, and each and every read request is satisfied by a random value of the legal input type. Write requests are also satisfied. We cojoin entry and return from read into a single read event, and entry and return from write into a single write event, since we are only interested in the behaviour of the fibre.

The sequence of all possible events which the fibre may exhibit is the coroutines control type. Note, the control type is a property of the coroutine (procedure).

8 Encoding Control Types

In general, the control type of a coroutine can be quite complex. However for special cases, a simple encoding can be given.

8.1 One shots

A one-shot routine is one that exhibits a bounded number of events before suiciding. The three most common one shots are:

Value: type W A coroutine which writes a single value to a channel and then exits.

Result: type R A coroutine which reads a single value to from channel and then exits.

Function: type RW A coroutine which reads one value from a channel, calculates an answer, writes that down a channel and then exits.

8.2 Continuous devices

A continuous coroutine is one which does not exit. It can therefore terminate only by starvation or blockage. The three most common kinds of such devices are

Source: type W+ Writes a continuous stream of values to a channel.

Sink: type R+ Reads a continuous stream of values from a channel.

Transducer Reads and writes.

Because the sequence of events is a stream, we may use convenient notations to describe control types. If possible, a regular expression will be used. Sometimes, a grammar will be required. In other cases there is no simple notation for the behaviour of a coroutine.

We will use postfix + for repetition.

8.3 Transducer Types

A transducer which read a value, write a value, then loops back and repeats is called a *functional transducer*, it may be given the type (RW)+.

In a functional language, a partial function has no natural encoding. There are two common solutions. The first is to return an option type, say `Some v`, if there is a result, or `None` if there is not. This solution involves modifying the

codomain. The other solution is to restrict the domain so that the subroutine is a function.

Coroutines, however, represent partial functions naturally. If a value is read for which there is no result, none is written! The type of a *partial function transducer* is therefore given by $((R+)W)^+$, in other words multiple reads may occur for each write. Note that two write may not occur in succession.

This type may also be applied to many other coroutines, for example the list filter higher order function.

8.4 Duality

Coroutines are dual to functions. The core difference is that they operate in time not space. Thus, in the dual space a spatial product type becomes a temporal sequence.

Coroutines are ideal for processing streams. Whereas function code cannot construct streams without laziness, and cannot deconstruct them without eagerness, coroutines are neither eager nor lazy.

One may view an eager functional application as driving a value into a function to get a result, and a lazy application as pulling a values into a function. Pushing value implies eagerly evaluating it, pulling implies the value is calculated on demand.

Coroutine simultaneously push and pull values across channels and so eliminate the evaluation model dichotomy that plagues functional programming. This coherence does not come for free: it is replaced by indeterminate event ordering.

9 Composition

By far the biggest advantage of coroutine modelling is the ultimate flexibility of composition. Coroutines provide far better modularity and reusability than functions, but this comes at the price of complexity. You will observe considerably more housekeeping is required to compose coroutines that procedures or functions, because, simply, there are more way to compose them.

A collection of coroutines can be regarded as black boxes resembling chips on a circuit board, with the wires connecting pins representing channels. So instead of using variables and binding constructions, we can construct more or less arbitrary networks.

9.1 Pipelines

The simplest kind of composition is the pipeline. It is a sequence of transducers wired together with the output of one transducer connected by a channel to the input of the next.

If the pipeline consists entirely of transducers is is an open pipeline. If there is a source at one end and a sink at the other it is a closed pipeline. Partially open pipelines can also exist.

The composition of two transducers has a type dependent on the left and right transducer types. If these types are functional, the result is functional. However for the more general case of two transducers of type $((R1+)W1+)$ and $((R2+)W2+)$ the result is somewhat suprising.

Connecting two transducers with a channel involves a theoretical abstraction operator which removes the write of the left transducer and the read of the right transducer from the behaviour of the pair: the combined transducer is has two unconnected channels, so it is still a two pin device. You may think the control type is $((R1+)W2+)$ but this is not the case!

The following behaviour is conforming: First we have left reading on R1, then writing W1. Right then reads R2. Now, left proceeds and reads R2, then writes W1. At this point right goes off and write W2, goes back, reads R2 again, goes off, and writes W2.

So what we have observed is: R1, R1, W2, W2 which is not recognised by $((R1+)W2)+$ because that requires R1 immediately after a W2. The problem is, each coroutine has memory, and acts as a buffer for at least a single value, so when two fibres are composed, there is one extra step of indeterminacy! The correct type is, in fact $((R1+)W2?W2)+$ where the question mark (?) indicates an optional operation.

It is clearer with a functional transducer, the composition of $(R1W1)+$ with $(R2W2)+$ is in fact $((R1?R1)(W2?W2))+$, meaning you can get two reads or two writes in a row. In fact this composite type is overly permissive, since the number of reads is at most two more than the number of writes at any time.

For a pipeline of N functional transducers, the output can be out of phase with the input by up to $N-1/2$ cycles. We put the $1/2$ there because there is always at least a half cycle delay between reading and writing, even with a tight loop. What happens depends on how the pipeline is driven.

If in the environment we write the left end then read the right end, the input and output are always coupled. However we can pump N values into the pipeline before reading. The reads out of the right end always correspond in sequence with the writes into the left end, but the coupling between reads and writes is sloppy.

This problem makes things much more difficult in more advanced topologies,

for example if two sources are driving a pipeline, there is no way to tell which input corresponds with which outputs. We cannot tell which source provided the data, even with a single functional transducer.

9.2 Clocks

Electronics designers know all about this problem and have a solution: clocking. If we add to each transducer an extra clock input which must be read prior to reading data, and add at the end a special feedback routine which writes the result first, then sends an acknowledge on the feedback channel, then we can construct a clock which reads the feedback, then writes a clock pulse to each transducer. This prevents the leftmost transducer reading any more input until the output is written. A one shot is used to bootstrap the process.

We can also organise things so that each transducer buffers a value, ensuring the pipeline is fully saturated.

9.3 Buffers

A transducer which simply copies its input to its output has no functional significance on a calculation, however it provides a one value buffer.

A more sophisticated device is called a *dynamic buffer*. For each input read by the buffer, a new fibre is spawned with a connecting channel which the input is written to. The new fibre is a one shot buffer which then writes the value and suicides. Such a buffer is equivalent to an unbounded set, where the word "set" implies a lack of ordering. All the dynamic buffers created try to write on the same channel so there is no assurance which will go first.

Such an unbounded set can, of course, grow forever unless the environment ensures this cannot happen. In a parser, for example, a dynamic buffer can be used for recursion in the same way a functional program uses the machine stack. Although both can overflow, both are bounded by the size and structure of the input being parsed.

10 Felix Implementation

The following functions and procedures are provided in Felix:

```
spawn_fthread: 1 -> 0
mk_ioschannel_pair[T]: 1 -> ischannel[T] * oschannel[T];
read[T]: ischannel[T] -> T
write[T]: oschannel[T] * T -> 0
```

In the abstract, channels are bidirectional and untyped. However we will restrict our attention to channels typed to support either read (`ischannel`) or write (`oschannel`) of a value of a fixed data type.

The following shorthand types are available:

```
%<T    ischannel[T]
%>T    oschannel[T]
```

More advanced typing exploiting channel capabilities are discussed later.

Simple example program:

```
proc demo () {
  var inp, out = mk_ioschannel_pair[int]();

  proc source () {
    for i in 1..10 perform write (out,i);
  }

  proc sink () {
    while true do
      var j = read inp;
      println$ j;
    done
  }

  spawn_fthread source;
  spawn_fthread sink;
}
demo();
```

In this program, we create a channel with an input and output end typed to transfer an `int`. The source coroutine writes the integers from 1 through to 10 inclusive to the write end of the channel, the sink coroutine reads integers from the channel and prints them.

The main fibre calls the `demo` procedure which launches two fibres with initial continuations the closures of the source and sink procedures.

When `demo` returns, the main fibre's current continuation no longer knows the channel, so the channel is not reachable from the main fibre.

The source coroutine returns after sending 10 integers to the sink via the channel. When a fibre no longer has a current continuation, returning to the non-existent caller causes the fibre to no longer have a legal state. This is known as suicide.

After the sink has read the last value, it becomes permanently Hungry. The sink procedure dies by starvation.

All fibres which die do so either by suicide, starvation, or blockage. Dead fibres will be reaped by the garbage collector provided they're unreachable. It is important for the creator of fibres and their connecting channels to forget the channels to ensure this occurs.

Unlike typical pre-emptive threading systems, deadlock is not an error. However a lock up which should lead to reaping of fibres but which fails to do so because they remain reachable is universally an error. This is known as a livelock: it leads to zombie fibres.

This usually occurs because some other fibre is statically capable of resolving the lockup, but does not do so dynamically. To prevent livelocks, variables holding channel values to which no I/O will occur dynamically should also go out of scope.